

NETAJI SUBHAS UNIVERSITY OF TECHNOLOGY



AI Hardware and tools

Lab File – VI Semester

TASK-1

Submitted by –

Paras Gunjyal (2021UCA1849)

Abhi Kansal (2021UCA1863)

Shivam Gupta (2021UCA1816)

CSAI – 1

1 Explore Basic Data Structure in R.

In R, there are several basic data structures that are commonly used for storing and manipulating data. Here are some of the fundamental data structures in R:

Vectors:

A vector is the most basic data structure in R, representing a one-dimensional array of elements. Elements in a vector must be of the same data type. You can create a vector using the `c()` function.

```
In [1]: numeric_vector <- c(12, 21, 14, 4, 5)

character_vector <- c("car", "bike", "bus")
```

```
In [2]: print(numeric_vector)
print(character_vector)

[1] 12 21 14 4 5
[1] "car" "bike" "bus"
```

Matrices:

A matrix is a two-dimensional data structure in R. It is created using the `matrix()` function.

```
In [3]: my_matrix <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2, ncol = 3)
print(my_matrix)

      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
```

Arrays:

An array is a multi-dimensional extension of a matrix. You can create an array using the `array()` function.

```
In [4]: # Creating a 3-dimensional array
my_array <- array(c(1:13), dim = c(2, 3, 2))
print(my_array)

, , 1

      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6

, , 2

      [,1] [,2] [,3]
[1,]    7    9   11
[2,]    8   10   12
```

Lists:

A list is a versatile data structure that can store elements of different data types. You can create a list using the `list()` function.

```
In [5]: my_list <- list(name = "Paras", age = 25, CCS_score = c(90, 85, 92))
print(my_list)
```

```
$name
[1] "Paras"
```

```
$age
[1] 25
```

```
$CCS_score
[1] 90 85 92
```

Data Frames:

A data frame is a two-dimensional table where each column can be of a different data type. It is created using the `data.frame()` function.

```
In [6]: my_data_frame <- data.frame(name = c("Ram", "Raju", "Charlie"),
                                   age = c(28, 22, 35),
                                   score = c(95, 89, 75))
print(my_data_frame)
```

```
  name age score
1   Ram  28   95
2  Raju  22   89
3 Charlie 35   75
```

Factors:

Factors are used to represent categorical data in R. They are created using the `factor()` function.

```
In [7]: my_factor <- factor(c("low", "medium", "high", "low", "medium"))
print(my_factor)
```

```
[1] low    medium high    low    medium
Levels: high low medium
```

2. Implement Linear Regression in R and Visualize the results.

```
In [8]: # Read the CSV file
insurance_data <- read.csv("insurance.csv")

# View the structure of the dataset
str(insurance_data)
```

```
'data.frame':  1338 obs. of  7 variables:
 $ age      : int  19 18 28 33 32 31 46 37 37 60 ...
 $ sex      : Factor w/ 2 levels "female","male": 1 2 2 2 2 1 1 1 2 1 ...
 $ bmi      : num  27.9 33.8 33 22.7 28.9 25.7 33.4 27.7 29.8 25.8 ...
 $ children: int   0 1 3 0 0 0 1 3 2 0 ...
 $ smoker   : Factor w/ 2 levels "no","yes": 2 1 1 1 1 1 1 1 1 1 ...
 $ region   : Factor w/ 4 levels "northeast","northwest",...: 4 3 3 2 2 3 3 2 1 2 ...
 $ expenses: num  16885 1726 4449 21984 3867 ...
```

```
In [9]: # Fit a linear regression model for predicting expenses
lm_model <- lm(expenses ~ age + bmi + children + smoker + region, data = insurance_data)

# Summary of the linear regression model
summary(lm_model)
```

Call:
lm(formula = expenses ~ age + bmi + children + smoker + region,
data = insurance_data)

Residuals:

Min	1Q	Median	3Q	Max
-11365.0	-2839.4	-985.3	1375.5	29924.5

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-11993.31	978.75	-12.254	< 2e-16	***
age	256.96	11.89	21.609	< 2e-16	***
bmi	338.76	28.56	11.862	< 2e-16	***
children	474.75	137.74	3.447	0.000585	***
smokeryes	23835.24	411.84	57.875	< 2e-16	***
regionnorthwest	-352.01	476.11	-0.739	0.459825	
regionsoutheast	-1034.93	478.53	-2.163	0.030738	*
regionsouthwest	-958.63	477.76	-2.007	0.045003	*

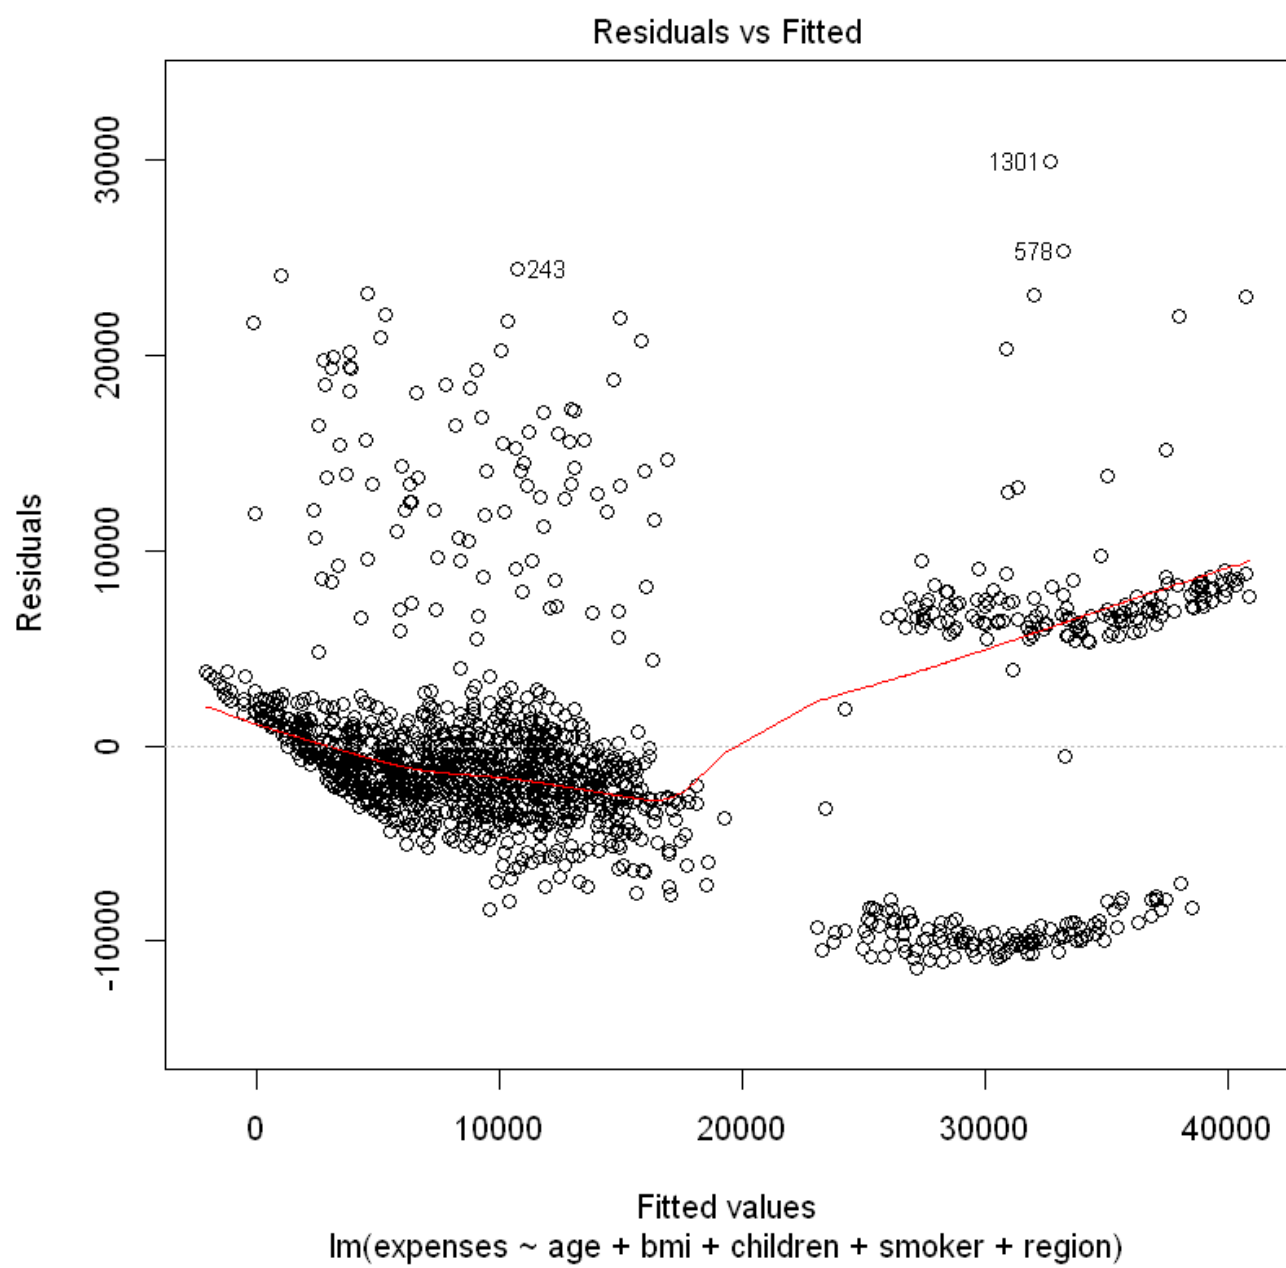
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 6060 on 1330 degrees of freedom
Multiple R-squared: 0.7509, Adjusted R-squared: 0.7496
F-statistic: 572.7 on 7 and 1330 DF, p-value: < 2.2e-16


```
In [10]: anova(lm_model)
```


	Df	Sum Sq	Mean Sq	F value	Pr(>F)
age	1	17530192069	17530192069	477.356520	1.135857e-90
bmi	1	5460652936	5460652936	148.696505	1.688995e-32
children	1	571896996	571896996	15.573062	8.352503e-05
smoker	1	123436032578	123436032578	3361.229288	0.000000e+00
region	3	233220248	77740083	2.116904	9.627598e-02
Residuals	1330	48842226838	36723479	NA	NA

```
In [11]: plot(lm_model)
```









3. Implement Logistic Regression in R and Visualize the results.

```
In [12]: # Read the CSV file
heart_data <- read.csv("heart.csv")

# View the structure of the dataset
str(heart_data)
```



```
'data.frame': 303 obs. of 14 variables:
 $ age      : int  63 37 41 56 57 57 56 44 52 57 ...
 $ sex      : int  1 1 0 1 0 1 0 1 1 1 ...
 $ cp       : int  3 2 1 1 0 0 1 1 2 2 ...
 $ trestbps : int  145 130 130 120 120 140 140 120 172 150 ...
 $ chol     : int  233 250 204 236 354 192 294 263 199 168 ...
 $ fbs      : int  1 0 0 0 0 0 0 0 1 0 ...
 $ restecg  : int  0 1 0 1 1 1 0 1 1 1 ...
 $ thalach  : int  150 187 172 178 163 148 153 173 162 174 ...
 $ exang     : int  0 0 0 0 1 0 0 0 0 0 ...
 $ oldpeak  : num  2.3 3.5 1.4 0.8 0.6 0.4 1.3 0 0.5 1.6 ...
 $ slope    : int  0 0 2 2 2 1 1 2 2 2 ...
 $ ca       : int  0 0 0 0 0 0 0 0 0 0 ...
 $ thal     : int  1 2 2 2 2 1 2 3 3 2 ...
 $ target   : int  1 1 1 1 1 1 1 1 1 1 ...
```

```
In [13]: # Fit Logistic regression model
logit_model <- glm(target ~ age + sex + cp + trestbps + chol + fbs + restecg + thalach + exang,
                  family = "binomial")

# Summary of the Logistic regression model
summary(logit_model)
```

Call:

```
glm(formula = target ~ age + sex + cp + trestbps + chol + fbs +
    restecg + thalach + exang + oldpeak + slope + ca + thal,
    family = "binomial", data = heart_data)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-2.5849  -0.3872   0.1551   0.5863   2.6249
```

Coefficients:

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  3.450472    2.571479   1.342 0.179653
age          -0.004908    0.023175  -0.212 0.832266
sex          -1.758181    0.468774  -3.751 0.000176 ***
cp           0.859851    0.185397   4.638 3.52e-06 ***
trestbps     -0.019477    0.010339  -1.884 0.059582 .
chol         -0.004630    0.003782  -1.224 0.220873
fbs          0.034888    0.529465   0.066 0.947464
restecg      0.466282    0.348269   1.339 0.180618
thalach      0.023211    0.010460   2.219 0.026485 *
exang        -0.979981    0.409784  -2.391 0.016782 *
oldpeak      -0.540274    0.213849  -2.526 0.011523 *
slope        0.579288    0.349807   1.656 0.097717 .
ca           -0.773349    0.190885  -4.051 5.09e-05 ***
thal         -0.900432    0.290098  -3.104 0.001910 **
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 417.64  on 302  degrees of freedom
Residual deviance: 211.44  on 289  degrees of freedom
AIC: 239.44
```

Number of Fisher Scoring iterations: 6

```
In [16]: # Visualize the results
# Let's create a plot for the probability of having heart disease based on age
new_data <- data.frame(age = seq(min(heart_data$age), max(heart_data$age), length.out = 100))
new_data$target_prob <- predict(logit_model, newdata = new_data, type = "response")
```

```
Error in eval(predvars, data, env): object 'sex' not found
```

```
Traceback:
```

```
1. predict(logit_model, newdata = new_data, type = "response")
2. predict.glm(logit_model, newdata = new_data, type = "response")
3. predict.lm(object, newdata, se.fit, scale = 1, type = if (type ==
.   "link") "response" else type, terms = terms, na.action = na.action)
4. model.frame(Terms, newdata, na.action = na.action, xlev = object$xlevels)
5. model.frame.default(Terms, newdata, na.action = na.action, xlev = object$xlevels)
6. eval(predvars, data, env)
7. eval(predvars, data, env)
```


```
In [15]: colnames(heart_data)
```

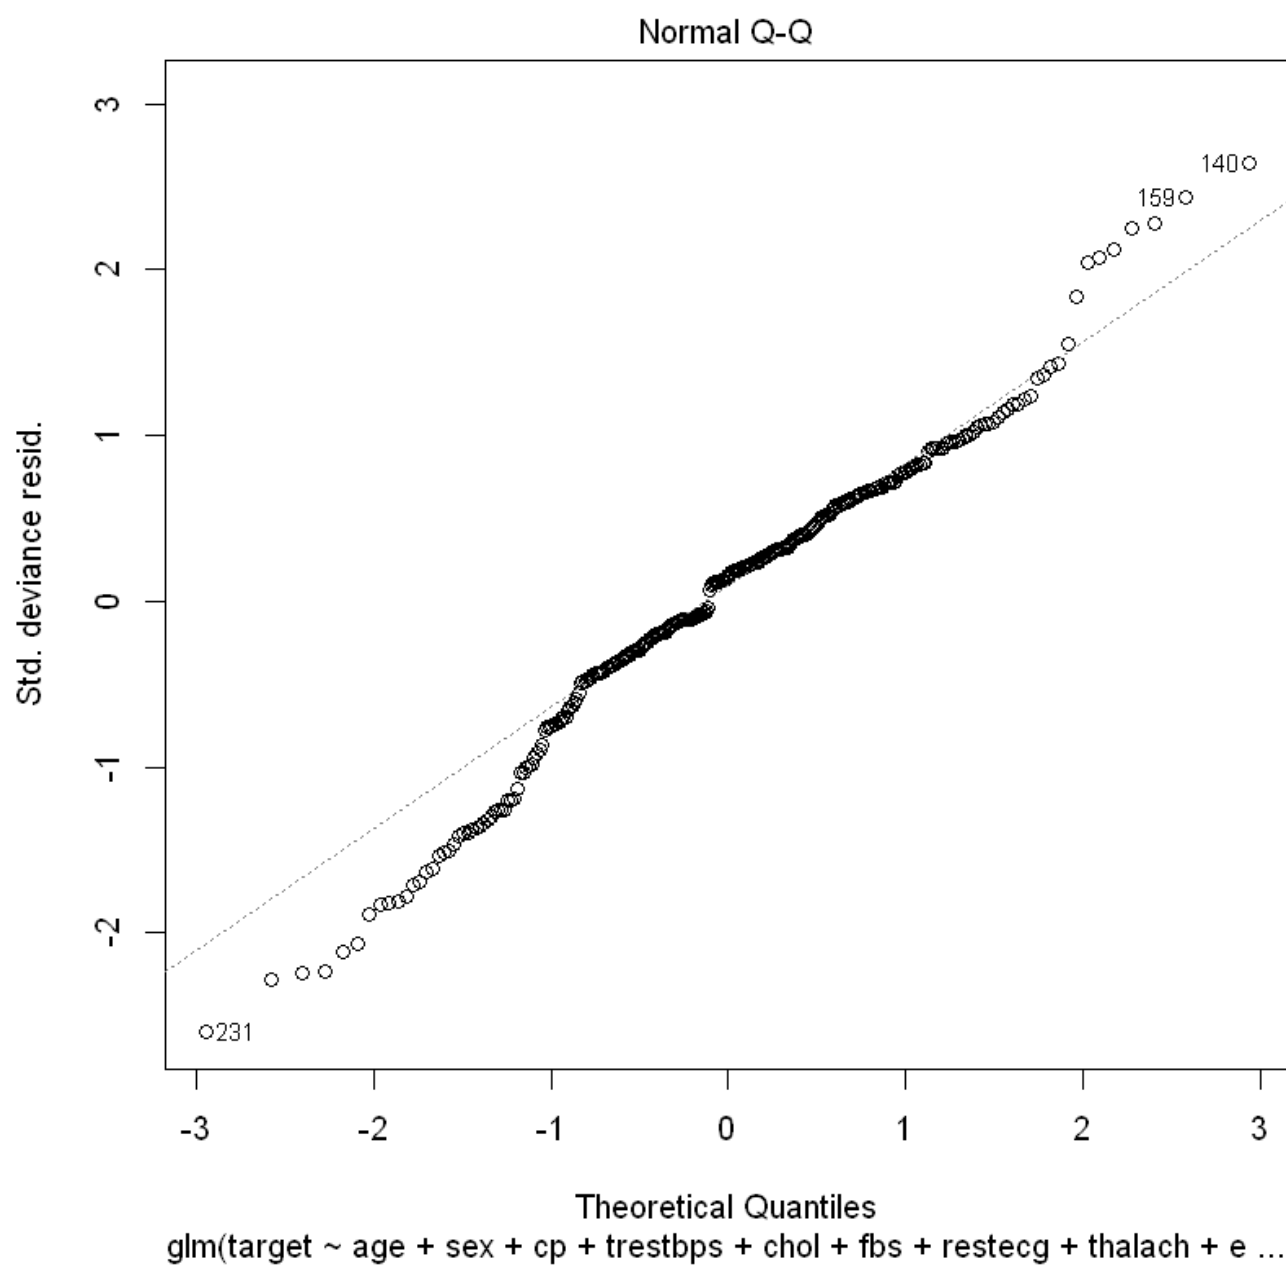
```
1. 'age'
2. 'sex'
3. 'cp'
4. 'trestbps'
5. 'chol'
6. 'fbs'
7. 'restecg'
8. 'thalach'
9. 'exang'
10. 'oldpeak'
11. 'slope'
12. 'ca'
13. 'thal'
14. 'target'
```

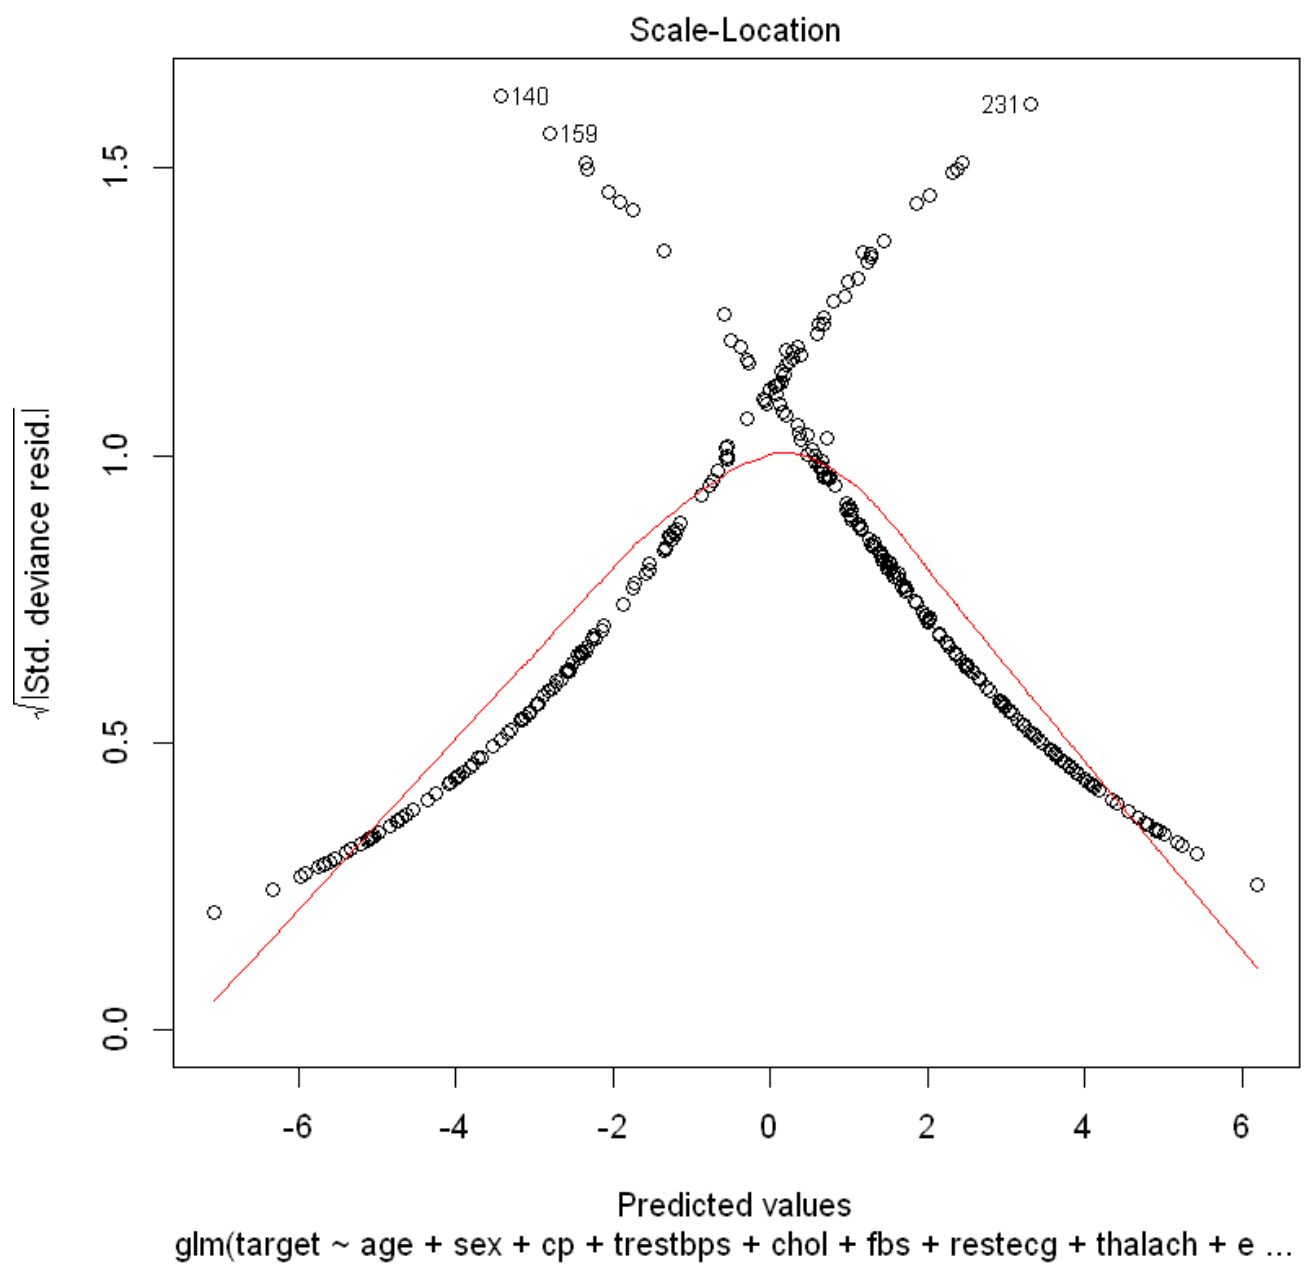
```
In [20]: anova(logit_model)
```


	Df	Deviance	Resid. Df	Resid. Dev
NULL	NA	NA	302	417.6381
age	1	15.7766919	301	401.8614
sex	1	31.2872127	300	370.5742
cp	1	59.7682306	299	310.8059
trestbps	1	6.7281625	298	304.0778
chol	1	1.9478502	297	302.1299
fbs	1	0.1098086	296	302.0201
restecg	1	1.7155609	295	300.3045
thalach	1	27.4338356	294	272.8707
exang	1	7.9650327	293	264.9057
oldpeak	1	22.2413718	292	242.6643
slope	1	1.0720758	291	241.5922
ca	1	20.2312203	290	221.3610
thal	1	9.9250308	289	211.4360

```
In [21]: plot(logit_model)
```









```
In [22]: # Fit a simpler logistic regression model (null model)
null_model <- glm(target ~ 1, data = heart_data, family = "binomial")

# Fit the full logistic regression model
full_model <- glm(target ~ age + sex + cp + trestbps + chol + fbs + restecg + thalach + exang

# Compare the models using anova
anova_result <- anova(null_model, full_model, test = "Chi")
print(anova_result)
```

Analysis of Deviance Table

Model 1: target ~ 1

Model 2: target ~ age + sex + cp + trestbps + chol + fbs + restecg + thalach +
exang + oldpeak + slope + ca + thal

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	302	417.64			
2	289	211.44	13	206.2	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

4. Implement any Machine learning Algorithm along with feature selection and data visualization on any dataset of your choice.

In [36]: *# Load necessary Libraries*

```
library(randomForest)
library(caret)
library(ggplot2)
```

In [39]: `titanic <- read.csv("Titanic.csv")`

In [43]: `titanic$Survived <- factor(titanic$Survived)`
`titanic$Age[is.na(titanic$Age)] <- mean(titanic$Age())`
`titanic$Fare[is.na(titanic$Fare)] <- mean(titanic$Fare)`
`titanic$Embarked[is.na(titanic$Embarked)] <- "Unknown"`

In [29]: *# Fit a Random Forest model*

```
rf_model <- randomForest(Species ~ ., data = train_data, ntree = 500)
```

```
# Load required libraries
library(randomForest)
library(caret)
library(ggplot2)

# Load the Titanic dataset
titanic <- read.csv("titanic.csv")

# Explore the dataset
head(titanic)
summary(titanic)

# Convert Survived to factor
titanic$Survived <- factor(titanic$Survived)

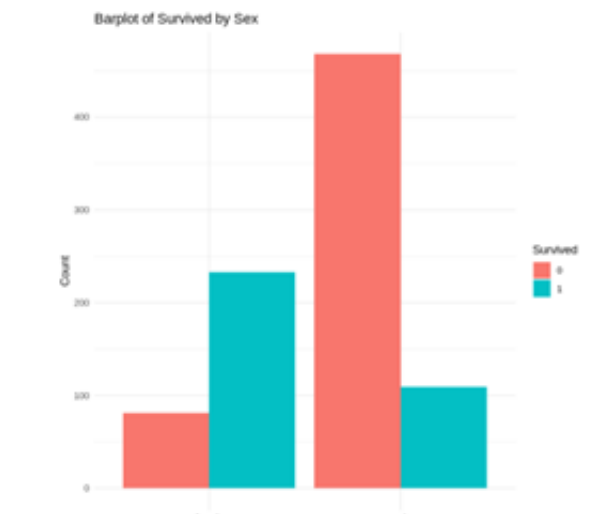
# Handle missing values
titanic$Age[is.na(titanic$Age)] <- mean(titanic$Age, na.rm = TRUE)
titanic$Fare[is.na(titanic$Fare)] <- mean(titanic$Fare, na.rm = TRUE)
titanic$Embarked[is.na(titanic$Embarked)] <- "Unknown"

# Data Visualization
# Barplot of Survived by Sex
ggplot(titanic, aes(x = Sex, fill = Survived)) +
  geom_bar(position = "dodge") +
  labs(x = "Sex", y = "Count", fill = "Survived", title = "Barplot of Survived by Sex") +
  theme_minimal()


# Barplot of Survived by Pclass
ggplot(titanic, aes(x = factor(Pclass), fill = Survived)) +
  geom_bar(position = "dodge") +
  labs(x = "Pclass", y = "Count", fill = "Survived", title = "Barplot of Survived by Pclass") +
  theme_minimal()

# Perform feature selection using caret package (wrapper method)
set.seed(123)
ctrl <- rfeControl(functions = rfFuncs, method = "cv", number = 10)
```


Sex	Age	SibSp	Parch
Length:891	Min. : 0.42	Min. :0.000	Min. :0.0000
Class :character	1st Qu.:20.12	1st Qu.:0.000	1st Qu.:0.0000
Mode :character	Median :28.00	Median :0.000	Median :0.0000
	Mean :29.70	Mean :0.523	Mean :0.3816
	3rd Qu.:38.00	3rd Qu.:1.000	3rd Qu.:0.0000
	Max. :80.00	Max. :8.000	Max. :6.0000
	NA's :177		
Ticket	Fare	Cabin	Embarked
Length:891	Min. : 0.00	Length:891	Length:891
Class :character	1st Qu.: 7.91	Class :character	Class :character
Mode :character	Median :14.45	Mode :character	Mode :character
	Mean :32.20		
	3rd Qu.:31.00		
	Max. :512.33		



'Survived'
1



Confusion Matrix

