

NETAJI SUBHAS UNIVERSITY OF TECHNOLOGY



AI Hardware and tools

Lab File – VI Semester

TASK-1

Submitted by –

Paras Gunjyal (2021UCA1849)

Abhi Kansal (2021UCA1863)

Shivam Gupta (2021UCA1816)

CSAI – 1

Emotion detector model and Tflite version for edge devices

The dataset is scrapped from google images and is then used for training the model

Importing Libraries

```
In [69]: 1
          2 import os
          3 import numpy as np
          4 import h5py
          5 import matplotlib.pyplot as plt
          6 import tensorflow as tf
          7 from tensorflow import keras
          8 from tensorflow.keras.models import load_model, Sequential
          9 from tensorflow.keras.layers import Flatten
         10 from tensorflow.keras.layers import Dense
         11 from tensorflow.keras.losses import SparseCategoricalCrossentropy
         12 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout
         13 from sklearn.metrics import accuracy_score
         14 from tensorflow.keras.metrics import Precision, Recall, BinaryAccuracy
         15 from sys import getsizeof
         16 import cv2
         17 import imghdr
         18 import pathlib
```

Defining Functions

```
In [2]: 1 def get_file_size(file_path):
          2     size = os.path.getsize(file_path)
          3     return size
          4
```

```
In [3]: 1 def convert_bytes(size, unit=None):
          2     if unit == "KB":
          3         return print('File size: ' + str(round(size / 1024, 3)) + ' Kilobytes')
          4     elif unit == "MB":
          5         return print('File size: ' + str(round(size / (1024 * 1024), 3)) + ' Megabytes')
          6     else:
          7         return print('File size: ' + str(size) + ' bytes')
          8
```

ML Model Creation

```
In [4]: 1 data_dir = 'data'
```

```
In [5]: 1 image_exts = ['.jpeg', '.jpg', '.bmp', '.png']
```

```
In [6]: 1 for image_class in os.listdir(data_dir):
2         for image in os.listdir(os.path.join(data_dir, image_class)):
3             image_path = os.path.join(data_dir, image_class, image)
4             try:
5                 img = cv2.imread(image_path)
6                 tip = imghdr.what(image_path)
7                 if tip not in image_exts:
8                     print('Image not in ext list {}'.format(image_path))
9                     os.remove(image_path)
10            except Exception as e:
11                print('Issue with image {}'.format(image_path))
12                # os.remove(image_path)
```

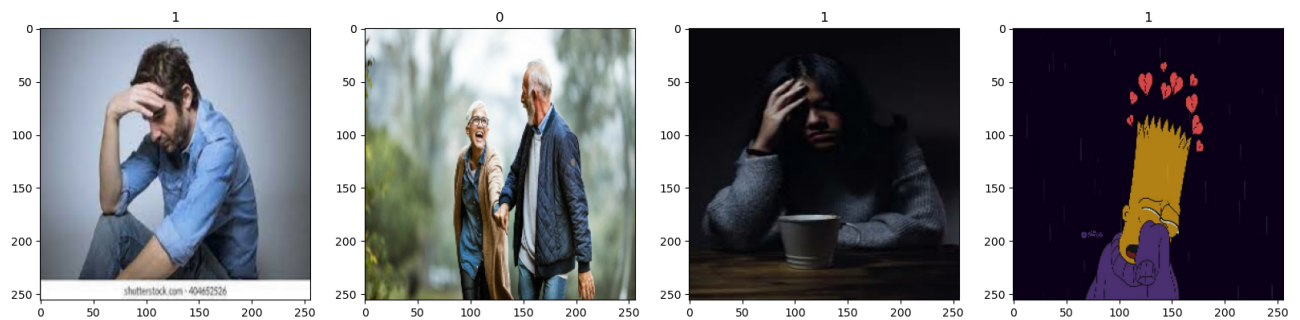
```
In [7]: 1 data = tf.keras.utils.image_dataset_from_directory('data')
```

Found 305 files belonging to 2 classes.

```
In [8]: 1 data_iterator = data.as_numpy_iterator()
```

```
In [9]: 1 batch = data_iterator.next()
```

```
In [10]: 1 fig, ax = plt.subplots(ncols=4, figsize=(20,20))
2         for idx, img in enumerate(batch[0][:4]):
3             ax[idx].imshow(img.astype(int))
4             ax[idx].title.set_text(batch[1][idx])
```



```
In [11]: 1 data = data.map(lambda x,y: (x/255, y))
```

```
In [12]: 1 data.as_numpy_iterator().next()
```

```
Out[12]: (array([[[[9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  ...,
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01]],
                  [[9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  ...,
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01]],
                  [[9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01],
                  [9.60784316e-01, 9.76470590e-01, 9.72549021e-01]]],
```

```
In [13]: 1 train_size = int(len(data)*.7)
          2 val_size = int(len(data)*.2)
          3 test_size = int(len(data)*.1)
```

```
In [14]: 1 train = data.take(train_size)
          2 val = data.skip(train_size).take(val_size)
          3 test = data.skip(train_size+val_size).take(test_size)
```

```
In [15]: 1 train
```

```
Out[15]: <TakeDataset element_spec=(TensorSpec(shape=(None, 256, 256, 3), dtype=tf.float32, name=None), TensorSpec(shape=(None,), dtype=tf.int32, name=None))>
```

```
In [16]: 1 model = Sequential()
```

```
In [17]: 1 model.add(Conv2D(16, (3,3), 1, activation='relu', input_shape=(256,256,3)))
          2 model.add(MaxPooling2D())
          3 model.add(Conv2D(32, (3,3), 1, activation='relu'))
          4 model.add(MaxPooling2D())
          5 model.add(Conv2D(16, (3,3), 1, activation='relu'))
          6 model.add(MaxPooling2D())
          7 model.add(Flatten())
          8 model.add(Dense(256, activation='relu'))
          9 model.add(Dense(1, activation='sigmoid'))
```

```
In [18]: 1 model.compile('adam', loss=tf.losses.BinaryCrossentropy(), metrics=['accuracy'])
```

In [19]:

1 model.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 254, 254, 16)	448
max_pooling2d (MaxPooling2D)	(None, 127, 127, 16)	0
conv2d_1 (Conv2D)	(None, 125, 125, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_2 (Conv2D)	(None, 60, 60, 16)	4624
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 16)	0
flatten (Flatten)	(None, 14400)	0
dense (Dense)	(None, 256)	3686656
dense_1 (Dense)	(None, 1)	257
=====		
Total params: 3,696,625		
Trainable params: 3,696,625		
Non-trainable params: 0		

In [20]:

1 logdir='logs'

In [21]:


1 tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=logdir)

In [22]: 1 hist = model.fit(train, epochs=20, validation_data=val, callbacks=[tensorboard_callback])

```
Epoch 1/20
7/7 [=====] - 6s 298ms/step - loss: 0.9341 - accuracy: 0.4688 - val_loss: 0.6971 - val_accuracy: 0.5156
Epoch 2/20
7/7 [=====] - 4s 330ms/step - loss: 0.6489 - accuracy: 0.5223 - val_loss: 0.6180 - val_accuracy: 0.5781
Epoch 3/20
7/7 [=====] - 3s 295ms/step - loss: 0.5666 - accuracy: 0.7143 - val_loss: 0.5027 - val_accuracy: 0.7969
Epoch 4/20
7/7 [=====] - 3s 302ms/step - loss: 0.4175 - accuracy: 0.8259 - val_loss: 0.3053 - val_accuracy: 0.8750
Epoch 5/20
7/7 [=====] - 3s 308ms/step - loss: 0.3012 - accuracy: 0.8884 - val_loss: 0.2213 - val_accuracy: 0.9531
Epoch 6/20
7/7 [=====] - 3s 298ms/step - loss: 0.2410 - accuracy: 0.9375 - val_loss: 0.1577 - val_accuracy: 0.9531
Epoch 7/20
7/7 [=====] - 3s 297ms/step - loss: 0.2175 - accuracy: 0.9241 - val_loss: 0.1422 - val_accuracy: 0.9844
Epoch 8/20
7/7 [=====] - 3s 305ms/step - loss: 0.1633 - accuracy: 0.9598 - val_loss: 0.1222 - val_accuracy: 0.9688
Epoch 9/20
7/7 [=====] - 3s 318ms/step - loss: 0.1080 - accuracy: 0.9688 - val_loss: 0.1504 - val_accuracy: 0.9531
Epoch 10/20
7/7 [=====] - 3s 313ms/step - loss: 0.1128 - accuracy: 0.9509 - val_loss: 0.1092 - val_accuracy: 0.9688
Epoch 11/20
7/7 [=====] - 3s 306ms/step - loss: 0.0665 - accuracy: 0.9777 - val_loss: 0.0566 - val_accuracy: 0.9844
Epoch 12/20
7/7 [=====] - 3s 297ms/step - loss: 0.0511 - accuracy: 0.9911 - val_loss: 0.0497 - val_accuracy: 0.9844
Epoch 13/20
7/7 [=====] - 3s 292ms/step - loss: 0.0667 - accuracy: 0.9821 - val_loss: 0.1139 - val_accuracy: 0.9375
Epoch 14/20
7/7 [=====] - 3s 308ms/step - loss: 0.0494 - accuracy: 0.9866 - val_loss: 0.0777 - val_accuracy: 0.9688
Epoch 15/20
7/7 [=====] - 3s 297ms/step - loss: 0.0758 - accuracy: 0.9866 - val_loss: 0.0290 - val_accuracy: 1.0000
Epoch 16/20
7/7 [=====] - 3s 297ms/step - loss: 0.0309 - accuracy: 0.9955 - val_loss: 0.0240 - val_accuracy: 1.0000
Epoch 17/20
7/7 [=====] - 3s 300ms/step - loss: 0.0194 - accuracy: 0.9955 - val_loss: 0.0390 - val_accuracy: 0.9844
Epoch 18/20
7/7 [=====] - 3s 309ms/step - loss: 0.0152 - accuracy: 0.9955 - val_loss: 0.0175 - val_accuracy: 1.0000
Epoch 19/20
7/7 [=====] - 3s 321ms/step - loss: 0.0081 - accuracy: 1.0000 - val_loss: 0.0050 - val_accuracy: 1.0000
Epoch 20/20
7/7 [=====] - 3s 300ms/step - loss: 0.0133 - accuracy: 0.9955 - val_loss: 0.0018 - val_accuracy: 1.0000
```


In [23]:

```
1 fig = plt.figure()
2 plt.plot(hist.history['loss'], color='teal', label='loss')
3 plt.plot(hist.history['val_loss'], color='orange', label='val_loss')
4 fig.suptitle('Loss', fontsize=20)
5 plt.legend(loc="upper left")
6 plt.show()
```



In [24]:

```
1 fig = plt.figure()
2 plt.plot(hist.history['accuracy'], color='teal', label='accuracy')
3 plt.plot(hist.history['val_accuracy'], color='orange', label='val_accuracy')
4 fig.suptitle('Accuracy', fontsize=20)
5 plt.legend(loc="upper left")
6 plt.show()
```



Testing the model

In [25]:

```
1 pre = Precision()
2 re = Recall()
3 acc = BinaryAccuracy()
4
5 for batch in test.as_numpy_iterator():
6     X, y = batch
7     yhat = model.predict(X)
8     pre.update_state(y, yhat)
9     re.update_state(y, yhat)
10    acc.update_state(y, yhat)
11
12 # Get the final results
13 precision_result = pre.result().numpy()
14 recall_result = re.result().numpy()
15 accuracy_result = acc.result().numpy()
16
17 # Print the results
18 print("Accuracy:", accuracy_result)
19 print("Precision:", precision_result)
20 print("Recall:", recall_result)
```

1/1 [=====] - 1s 521ms/step

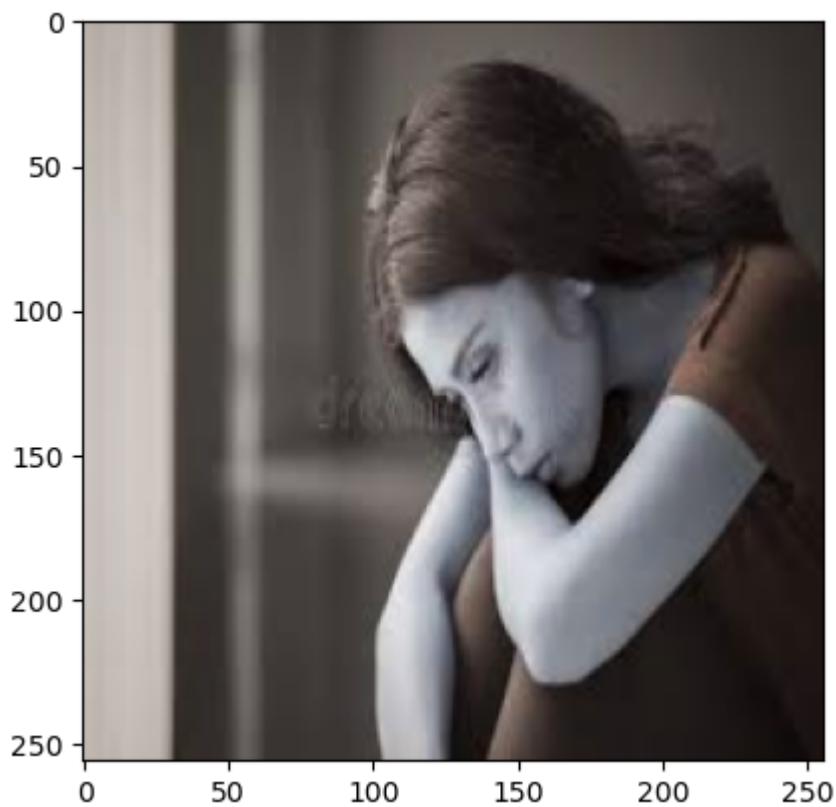
Accuracy: 1.0

Precision: 1.0

Recall: 1.0

In [26]:

```
1 # Load the image using OpenCV
2 img = cv2.imread('sad_test.jpg')
3
4 plt.imshow(img)
5 plt.show()
6
7 resize = tf.image.resize(img, (256,256))
8 plt.imshow(resize.numpy().astype(int))
9 plt.show()
10
```



```
In [27]: 1 result = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 102ms/step
```

```
In [28]: 1 result  
2
```



```
Out[28]: array([[0.9990477]], dtype=float32)
```

```
In [29]: 1 if result > 0.5:  
2     print(f'Predicted class is Sad')  
3 else:  
4     print(f'Predicted class is Happy')
```

```
Predicted class is Sad
```

In [30]:

```
1 # Load the image using OpenCV
2 img = cv2.imread('happy_test.jpg')
3
4 plt.imshow(img)
5 plt.show()
6
7 resize = tf.image.resize(img, (256,256))
8 plt.imshow(resize.numpy().astype(int))
9 plt.show()
10
```



```
In [31]: 1 result = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 17ms/step
```

```
In [32]: 1 result
```

```
Out[32]: array([[0.08620903]], dtype=float32)
```

```
In [33]: 1 if result > 0.5:
2         print(f'Predicted class is Sad')
3 else:
4         print(f'Predicted class is Happy')
```

```
Predicted class is Happy
```

```
In [34]: 1 model.save(os.path.join('models', 'SadUnsad.h5'))
```

```
In [35]: 1 new_model = load_model('models/SadUnsad.h5')
```

```
In [36]: 1 convert_bytes(get_file_size('models/SadUnsad.h5'), "MB")
```

```
File size: 42.355 Megabytes
```

Model size before any optimization

Tflite conversion

```
In [37]: 1 TF_LITE_MODEL_FILE_NAME = "tf_lite_model.tflite"
```

```
In [38]: 1 tf_lite_converter = tf.lite.TFLiteConverter.from_keras_model(new_model)
2 tflite_model = tf_lite_converter.convert()
3
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_com
piled_convolution_op, _jit_compiled_convolution_op while saving (showing 3 of 3). Th
ese functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: C:\Users\Paras\AppData\Local\Temp\tmptyubqarq\ass
ets
```

```
INFO:tensorflow:Assets written to: C:\Users\Paras\AppData\Local\Temp\tmptyubqarq\ass
ets
```

```
In [39]: 1 tflite_model_name = TF_LITE_MODEL_FILE_NAME
2 open(tflite_model_name, "wb").write(tflite_model)
3
```

```
Out[39]: 14790588
```

```
In [40]: 1 convert_bytes(get_file_size(TF_LITE_MODEL_FILE_NAME), "MB")
```

```
File size: 14.105 Megabytes
```

SIZE AFTER OPTIMIZATION

In [50]:

```
1 interpreter = tf.lite.Interpreter(model_path = TF_LITE_MODEL_FILE_NAME)
2 input_details = interpreter.get_input_details()
3 output_details = interpreter.get_output_details()
4 print("Input Shape:", input_details[0]['shape'])
5 print("Input Type:", input_details[0]['dtype'])
6 print("Output Shape:", output_details[0]['shape'])
7 print("Output Type:", output_details[0]['dtype'])
```

```
Input Shape: [ 1 256 256  3]
Input Type: <class 'numpy.float32'>
Output Shape: [1 1]
Output Type: <class 'numpy.float32'>
```

In [55]:

```
1 interpreter.resize_tensor_input(input_details[0]['index'], (1000, 256, 256, 3))
2 interpreter.resize_tensor_input(output_details[0]['index'], (1000, 10))
3 interpreter.allocate_tensors()
4 input_details = interpreter.get_input_details()
5 output_details = interpreter.get_output_details()
6 print("Input Shape:", input_details[0]['shape'])
7 print("Input Type:", input_details[0]['dtype'])
8 print("Output Shape:", output_details[0]['shape'])
9 print("Output Type:", output_details[0]['dtype'])
10
```

```
Input Shape: [1000 256 256  3]
Input Type: <class 'numpy.float32'>
Output Shape: [1000  1]
Output Type: <class 'numpy.float32'>
```

In [56]:

```
1 pre = Precision()
2 re = Recall()
3 acc = BinaryAccuracy()
4
5 for batch in test.as_numpy_iterator():
6     X, y = batch
7     yhat = model.predict(X)
8     pre.update_state(y, yhat)
9     re.update_state(y, yhat)
10    acc.update_state(y, yhat)
11
12 # Get the final results
13 precision_result = pre.result().numpy()
14 recall_result = re.result().numpy()
15 accuracy_result = acc.result().numpy()
16
17 # Print the results
18 print("Accuracy:", accuracy_result)
19 print("Precision:", precision_result)
20 print("Recall:", recall_result)
```

```
1/1 [=====] - 0s 24ms/step
Accuracy: 0.9411765
Precision: 0.875
Recall: 1.0
```

In [73]:

```
1 for file in pathlib.Path('data').iterdir():  
2     print("File Path:", file.resolve())  
3     # Read and resize the image...
```

File Path: C:\Users\Paras\6 sem\Ai HT\TinyMl\data\happy

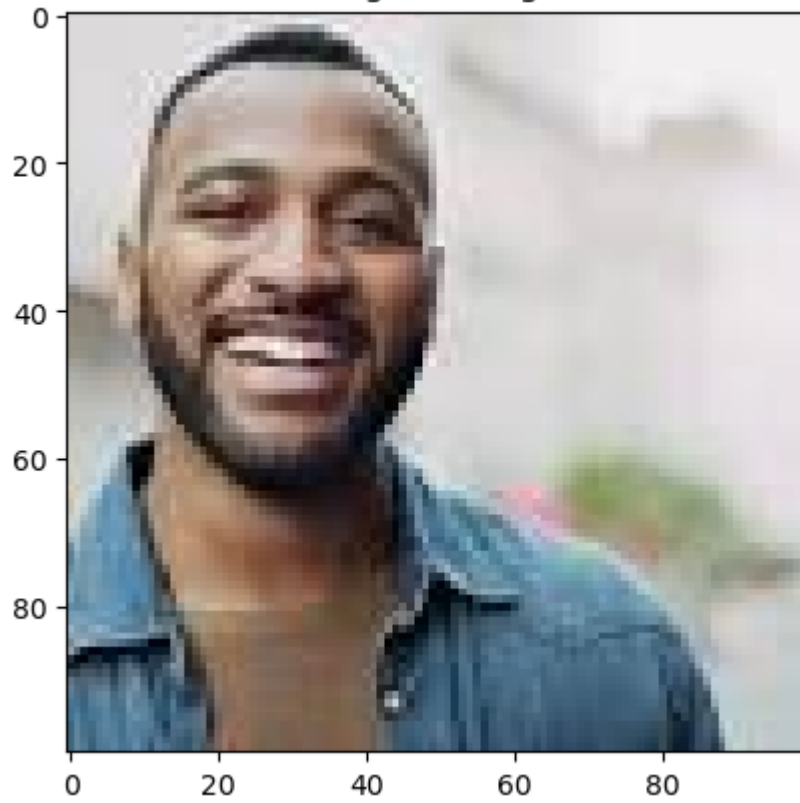
File Path: C:\Users\Paras\6 sem\Ai HT\TinyMl\data\sad

Testing the Tflite Model

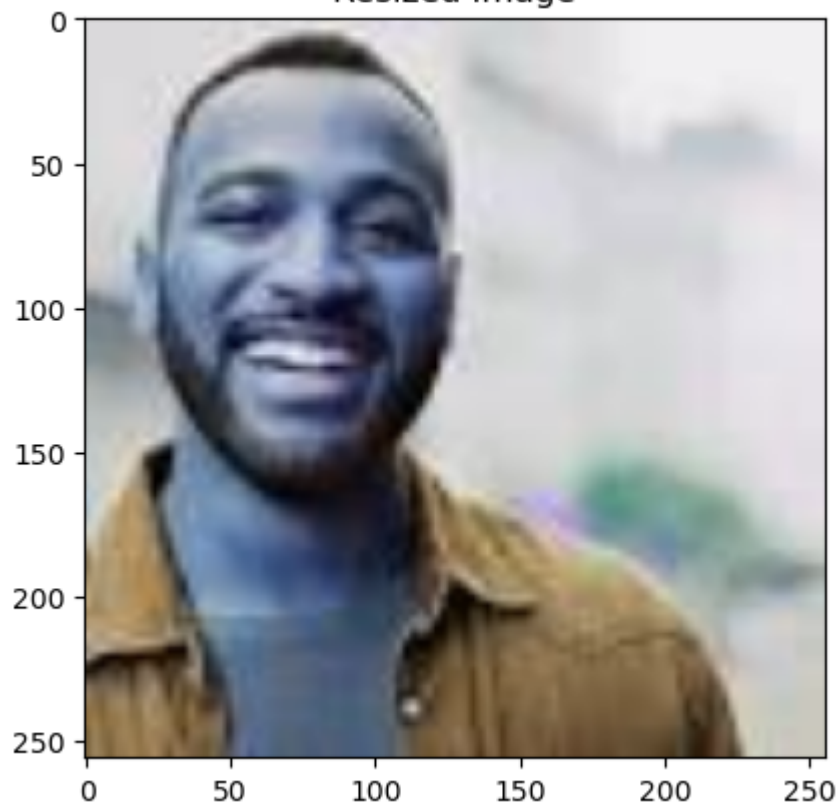
In [87]:

```
1
2 # Load TFLite model
3 model_path = 'tf_lite_model.tflite'
4 interpreter = tf.lite.Interpreter(model_path=model_path)
5 interpreter.allocate_tensors()
6
7 # Get input and output details
8 input_details = interpreter.get_input_details()
9 output_details = interpreter.get_output_details()
10
11 # Function to preprocess image
12 def preprocess_image(image_path):
13     image = Image.open(image_path)
14     image = image.resize((256, 256)) # Replace with the input size of your model
15     image = tf.keras.preprocessing.image.img_to_array(image)
16     image = tf.keras.applications.mobilenet_v2.preprocess_input(image)
17     return image
18
19 # Function to make predictions
20 def predict_image(image_path):
21     input_data = preprocess_image(image_path)
22     input_data = tf.expand_dims(input_data, axis=0)
23
24     interpreter.set_tensor(input_details[0]['index'], input_data)
25     interpreter.invoke()
26
27     output_data = interpreter.get_tensor(output_details[0]['index'])
28     return output_data
29
30 # Load image using OpenCV
31 img_path = 'happy_test.jpg'
32 img = cv2.imread(img_path)
33
34 # Display the original image
35 plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
36 plt.title("Original Image")
37 plt.show()
38
39 # Resize using TensorFlow
40 resize = tf.image.resize(tf.convert_to_tensor(img), (256, 256)).numpy().astype(int)
41
42 # Display the resized image
43 plt.imshow(resize)
44 plt.title("Resized Image")
45 plt.show()
46
47 # Predict using the TFLite model
48 prediction = predict_image(img_path)
49 predicted_class = int(prediction[0][0]) # Assuming a single output class, adjust
50
51 print(f"Prediction for {img_path}: Class {predicted_class}")
52
```


Original Image



Resized Image



Prediction for happy_test.jpg: Class 1

```
In [88]: 1 result = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 19ms/step
```

```
In [89]: 1 result
```

```
Out[89]: array([[0.08366012]], dtype=float32)
```

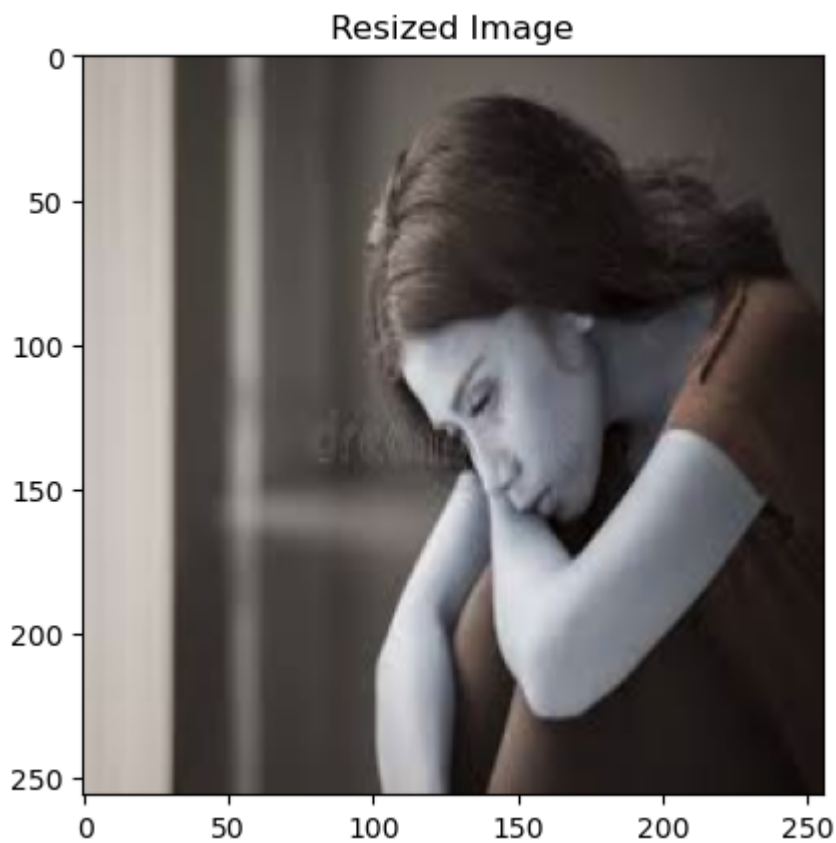
```
In [90]: 1 if result > 0.5:  
2     print(f'Predicted class is Sad')  
3 else:  
4     print(f'Predicted class is Happy')
```

Predicted class is Happy

In [91]:

```
1 img_path = 'sad_test.jpg'
2 img = cv2.imread(img_path)
3
4 # Display the original image
5 plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
6 plt.title("Original Image")
7 plt.show()
8
9 # Resize using TensorFlow
10 resize = tf.image.resize(tf.convert_to_tensor(img), (256, 256)).numpy().astype(int)
11
12 # Display the resized image
13 plt.imshow(resize)
14 plt.title("Resized Image")
15 plt.show()
16
17 # Predict using the TFLite model
18 prediction = predict_image(img_path)
19 predicted_class = int(prediction[0][0]) # Assuming a single output class, adjust accordingly
20
21 print(f"Prediction for {img_path}: Class {predicted_class}")
22
```





Prediction for sad_test.jpg: Class 1

```
In [92]: 1 result = model.predict(np.expand_dims(resize/255, 0))
```

```
1/1 [=====] - 0s 17ms/step
```

```
In [93]: 1 result
```

```
Out[93]: array([[0.99911946]], dtype=float32)
```

```
In [94]: 1 if result > 0.5:
2         print(f'Predicted class is Sad')
3 else:
4         print(f'Predicted class is Happy')
```

Predicted class is Sad

Conclusion

Hence, the Model is now finalized and deployable for edge devices for remote usage